

# A Deep Dive into Memory Access: Extended Roofline for Power Management

Presentation of Bachelor Thesis

First examiner: Prof. Dr. Matthias Müller

Second examiner: Prof. Dr. Martin Schulz

Adviser: M. Sc. Bo Wang

Lehrstuhl für Hochleistungsrechnen (Informatik 12), IT Center, RWTH Aachen

Lehrstuhl für Rechnerarchitektur & Parallele Systeme. Technische Universität München

Liangkun He, 27.07.2021

# Outline

---

- Motivation
- Background
- Roofline Model
- Extended Roofline Model
- Implementation
- Evaluation
- Summary

# Motivation

---

- Supercomputer's power consumption is significant
  - Fugaku is the top one supercomputer. It consumes 28.3MW to deliver 415.53 Pflops/s [1]
- Large power consumption brings many problems
  - High energy cost
  - Heats dissipation
  - High carbon dioxide emissions
- Exascale supercomputer ( $1\text{ExaFlop/s}=10^3\text{PFlop/s}$ )
  - Next generation supercomputers that yet to come
- Performance tuning
  - Power consumption cannot exceed the power budget
  - Performance and power consumption optimizations are important



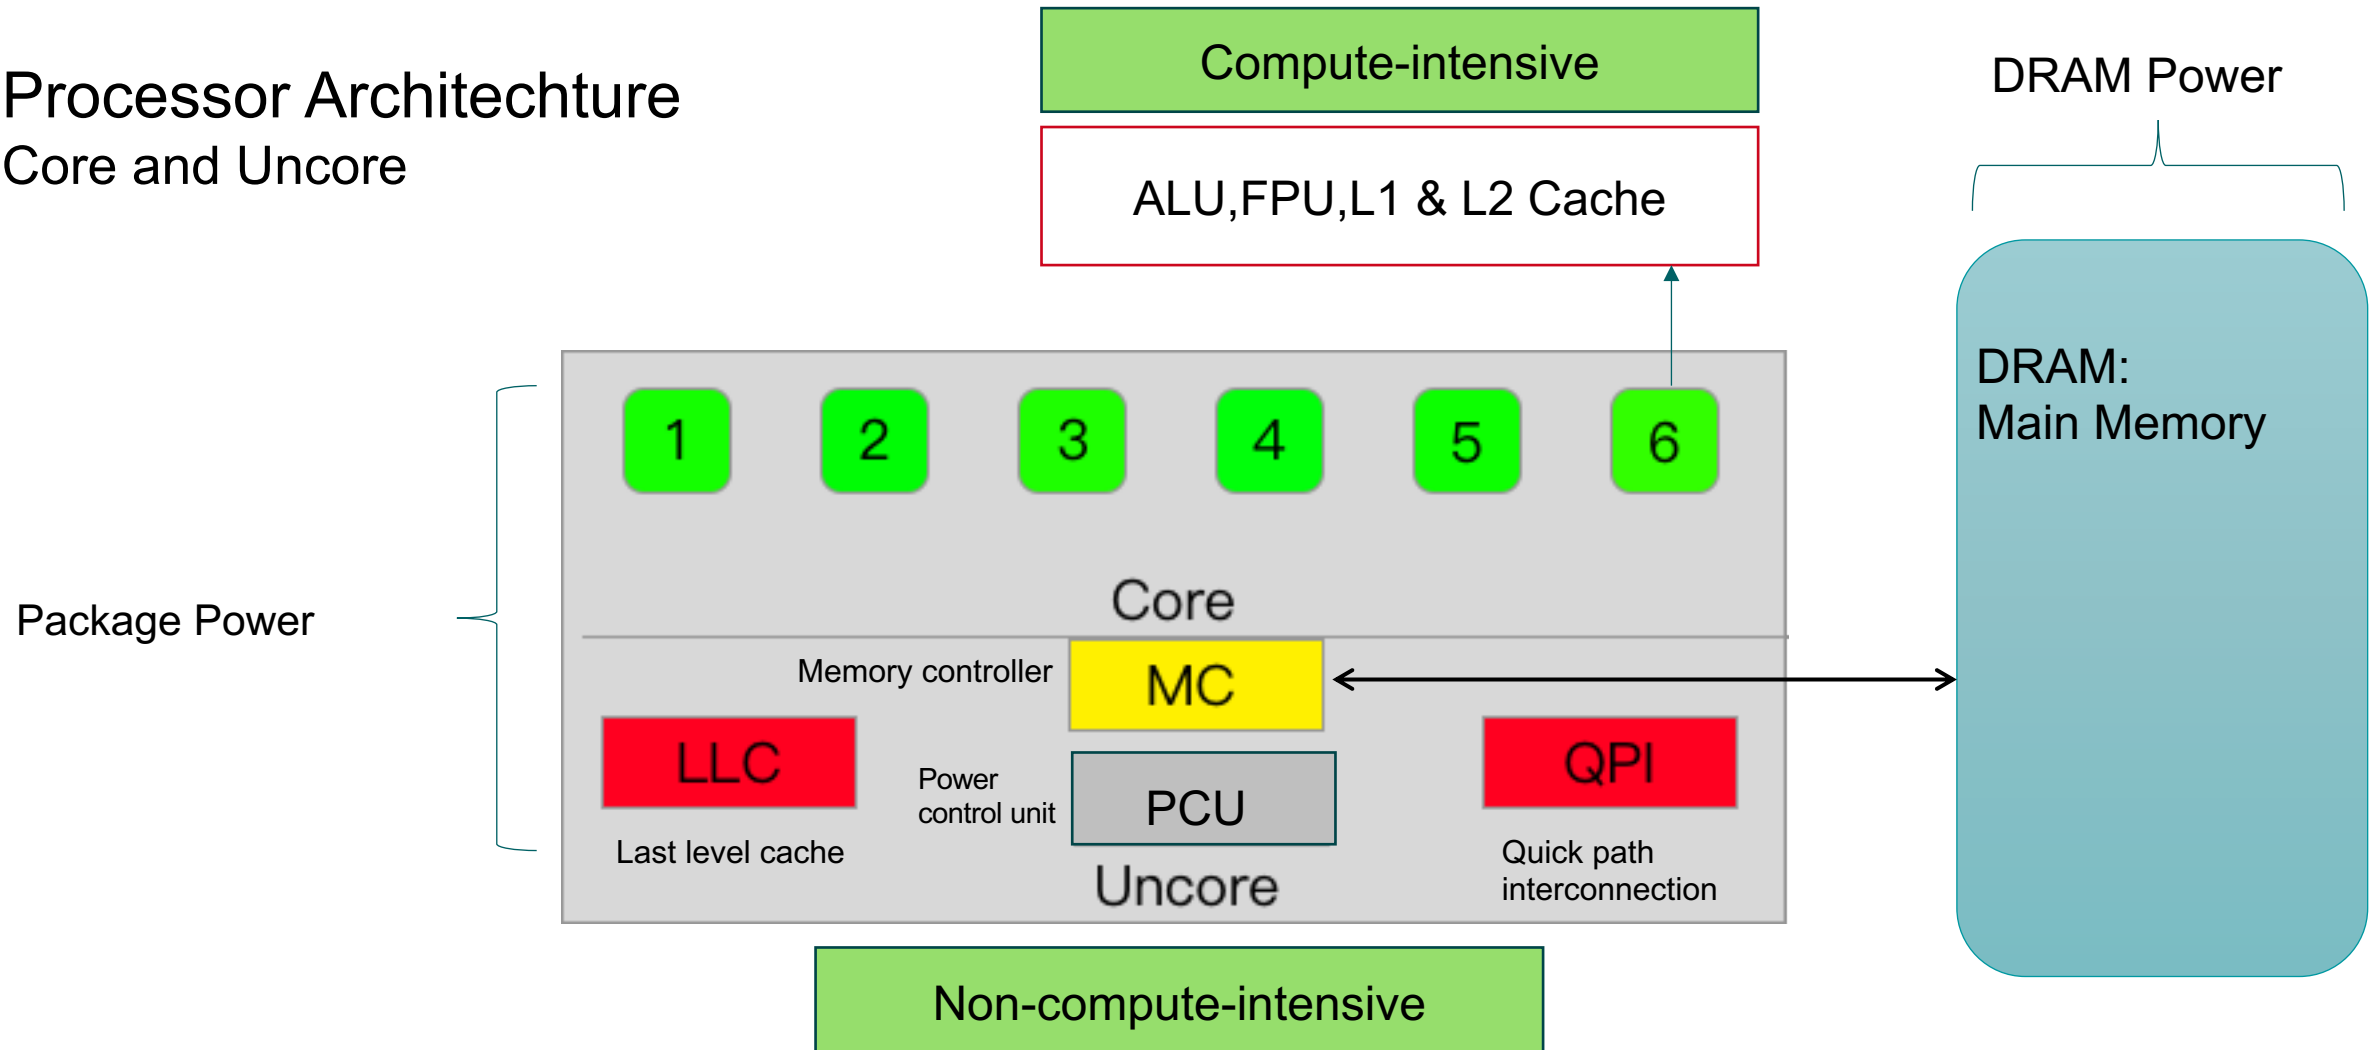
[2]

## Contribution

- Study the processor architecture
  - Core vs. uncore
  - Different components of a processor have their individual frequency
- Implement the Roofline model based on the **sampling** approach
  - Identify the execution patterns of an application for each sampling intervals
    - Compute bound or memory bound
  - Modify each component's frequency respectively during runtime to optimize the performance and power consumption
  - Propose the **time interval analysis** method
    - Analyse the activities of CPU and memory for each sampling intervals
    - Extend the Roofline model with the latency bottleneck
- Evaluate the implementation on different benchmarks

## Background

- Processor Architecture
- Core and Uncore




## Running Average Power Limit(RAPL)

- Set power cap for processors in a time interval
  - Keep the core or uncore frequency low to make sure they do not exceed the power cap
- RAPL is not aware of the **execution patterns** of the running application
  - **Compute intensive** or **memory intensive**
  - For memory intensive application such as Stream benchmark, RAPL may reduce the uncore frequency and thus lower the performance

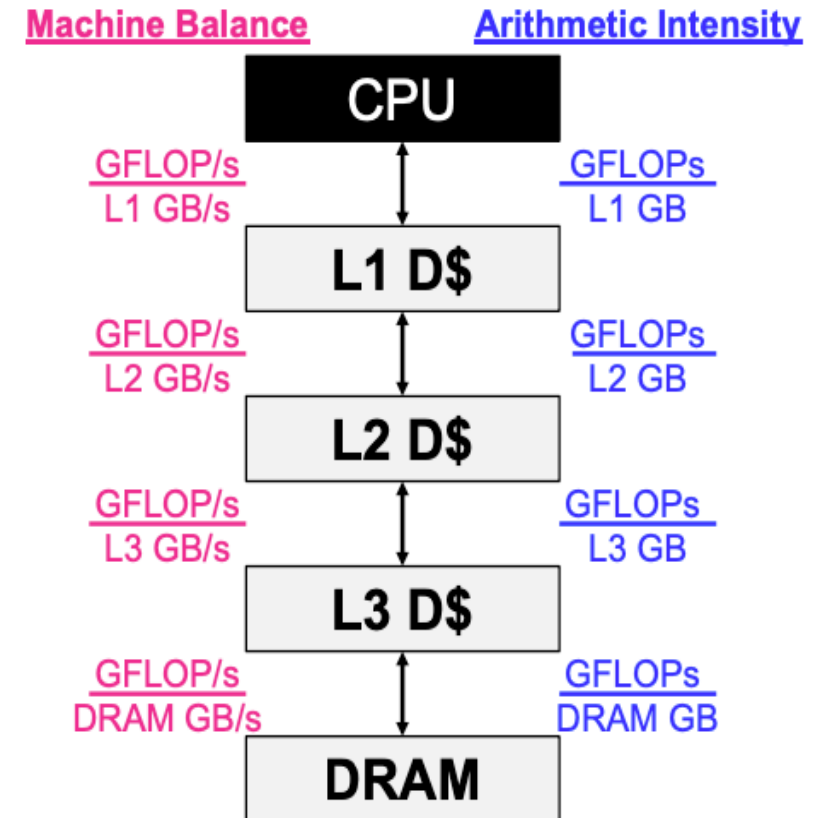
# Roofline Model

- $P = \min(P_{Peak}, I*B)$  [Ins/s]
  - Performance bottleneck is either
    - Computation:  $P_{Peak}$  [Ins/s]
    - Data transfer:  $I*B$  [Ins/Byte \* Byte/s]
- B: peak memory bandwidth
- I : arithmetic intensity [Ins/Byte]
  - Describe the number of instructions per data transfer
  - Using Ins/Byte has advantages over Flop/Byte in implementation
    - Applications may execute both integer and floating-point operations
    - Ins/Byte is easier to measure and more accurate
- Machine Balance
  - $P_{Peak} = I*B$
  - Best use of computational performance and memory bandwidth
- $I < \text{machine balance} \rightarrow \text{memory bound}$
- $I > \text{machine balance} \rightarrow \text{compute bound}$



## Memory hierarchy

- Cache
  - Small but fast memory that are integrated in the processor
- Each level has its own bandwidth hence different machine balances
- The amounts of data that transferred between memory levels are different hence different arithmetic intensities




[3]



# Roofline Model


## This work focuses on L3 cache and DRAM levels

- The performance is bound by the minimal attainable ins/s
  - $P = \min(P_{Peak}, I_{L3} * B_{L3}, I_{DRAM} * B_{DRAM})$
- Construct two bandwidth ceilings
  - L3 cache and DRAM
- Bound to L3 cache in the example




# Roofline Model

- Assumption
  - Computation and data transfer between CPU and memory overlap perfectly.(ignore memory latency)
- Fact
  - Memory latency may not be eliminated completely
- Out-of-order execution:

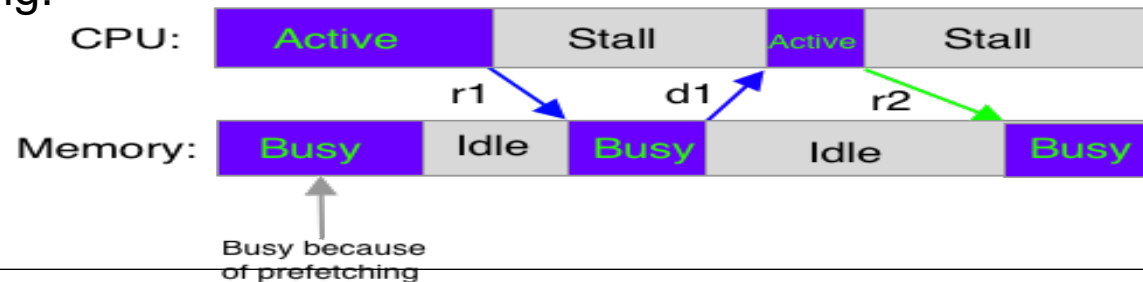


In-order execution




Out-of-order execution

- Hardware prefetching:



# Roofline Model


- How core and uncore frequency scaling affects the L3 and DRAM latency
- Latency is measured by tinyBench
- Latency increases when core/uncore frequency decreases



# Extended Roofline Model


## Time Interval Analysis

- Transfer delay: Part of memory latency that cannot be eliminated
- Time interval analysis
  - Calculate the transfer delay based on the activities of CPU and memory in a sampling interval
  - The summarized activities can be derived from hardware events




## Extended Roofline Model

- Transfer delay is approximated to  $\min(t_{Stall}^{cpu}, t_{Idle}^{mem})$ 
  - Time for synchronization between CPU and memory
  - When CPU stalls and memory is idle
- **Transfer delay bound** with memory domination
  - Slightly decrease the core frequency



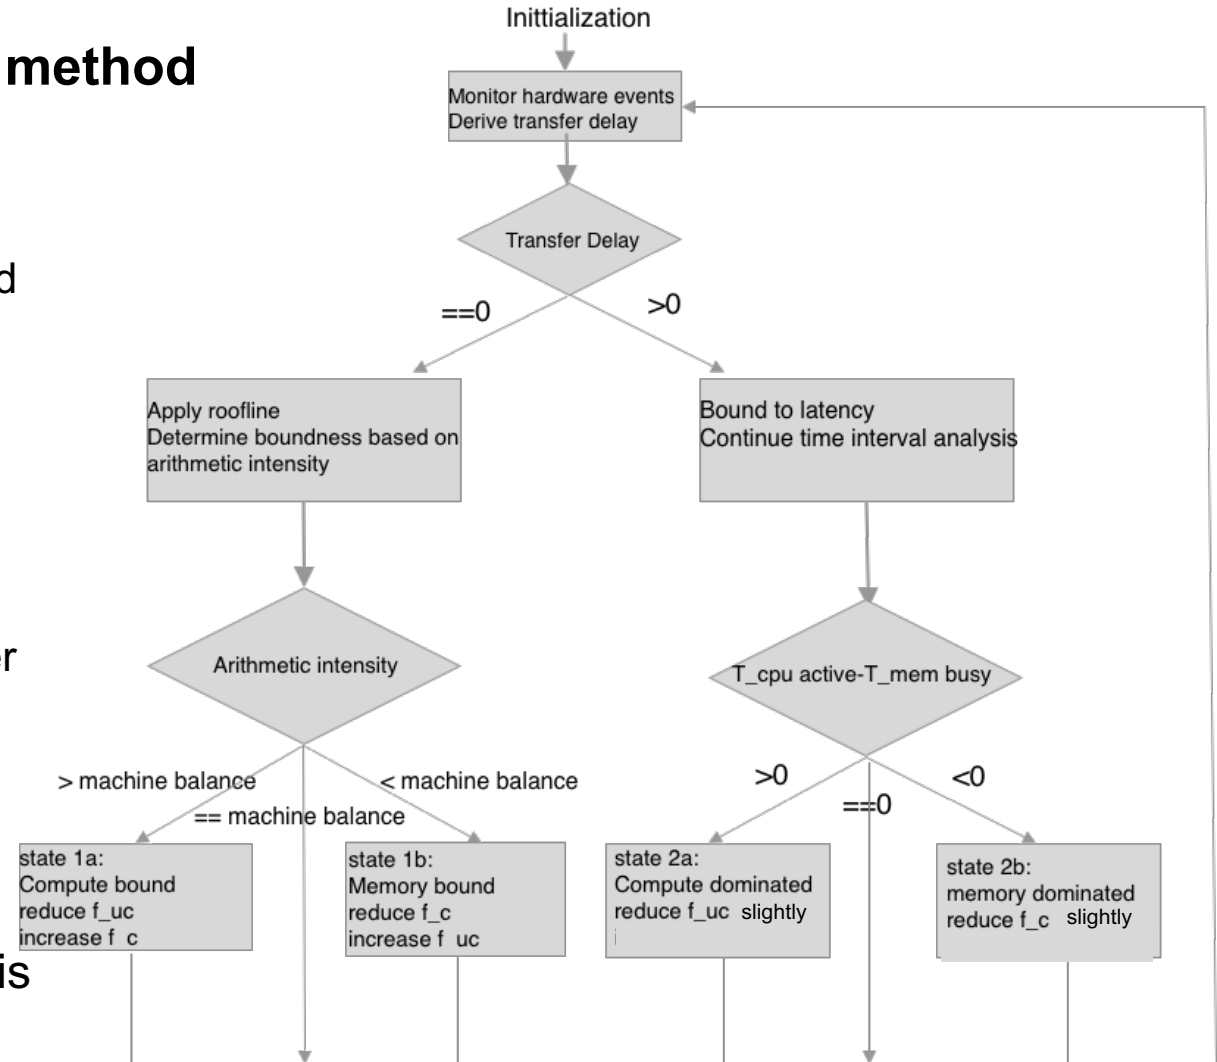
## Extended Roofline Model

- Transfer delay is approximated to  $\min(t_{Stall}^{cpu}, t_{Idle}^{mem})$
- Transfer delay bound with compute domination
  - Slightly decrease the uncore frequency



## Extended Roofline Model based on sampling method

- Sampling interval: 0.3s
  - Less over head
  - Detects some applications that vary from compute bound to memory bound.
- Initiation
  - Core and uncore frequency is set to maximum
- Measure hardware events
  - Core/uncore frequency, instructions, cycles, data transfer of LLC and DRAM, PKG energy and power...
  - Derive the transfer delay, CPU active and memory busy time, arithmetic intensities...
- If transfer delay approximates to zero → Roofline
- If transfer delay  $> 0$  → Continue time interval analysis



## Hardware Platform and Tools for Measuring Hardware Events

- Haswell Intel (R) Xeon (R) processor E5 v3 @ 2.3GHz
  - Core frequency: 1.2 GHz ~ 2.8 GHz
  - Uncore frequency: 1.2 GHz ~ 3.0 GHz
  - L3 cache: 45 MB
  - DRAM: 32 GB
- Perf
  - Measuring core events and DRAM read/write
  - Lack uncore support
- Libmsr
  - Measuring uncore activities: LLC read/write





## Benchmarks

- Stream (to evaluate the boundary cases)
  - Stream\_L3 : LLC intensive
    - Total memory required = 20.6 MiB < L3 cache size: suitable for L3 cache testing
  - Stream\_DRAM : DRAM intensive
    - Total memory required >> L3 cache size: suitable for DRAM testing
- NAS ***Parallel Benchmarks***
  - Memory intensive: Conjugate Gradient(CG) Multi-Grid (MG) Fourier Transform (FT)
  - Compute intensive: Embarrassingly Parallel (EP)
  - LLC intensive: SP
  - Different execution patterns: BT

# Evaluation


---

- Stream (without power cap):
  - Measure core and uncore frequency during runtime by the extended Roofline model
  - PKG power saving, energy saving and slowdown compared to the default ( $f_c = 2.8\text{GHz}$ ,  $f_{uc} = 3.0\text{GHz}$ )
- NAS Parallel Benchmarks:
  - Energy optimization ([without power cap](#)):
    - PKG power and energy saving, and slowdown compared to the default ( $f_c = 2.8\text{GHz}$ ,  $f_{uc} = 3.0\text{GHz}$ )
  - Performance and energy optimization ([under a power cap \(69W\)](#)):
    - Speedup and energy saving compared to the default




### PKG Power, energy saving, and slowdown [%] wrt. default

- For DRAM
  - up to 40% of PKG power or energy saving with less than 5% slowdown
- For L3 Cache
  - Marginal impacts on L3 cache



### BT's execution patterns (no power cap)

- Bounds to transfer delay but varies from compute domination to memory domination
- Core frequency: 2.3 ~ 2.8 GHz
- Uncore frequency : 2.2 ~ 3.0 GHz




### PKG power saving, slowdown and PKG energy saving [%] wrt. default (no power cap)

Up to 13% PKG power saving

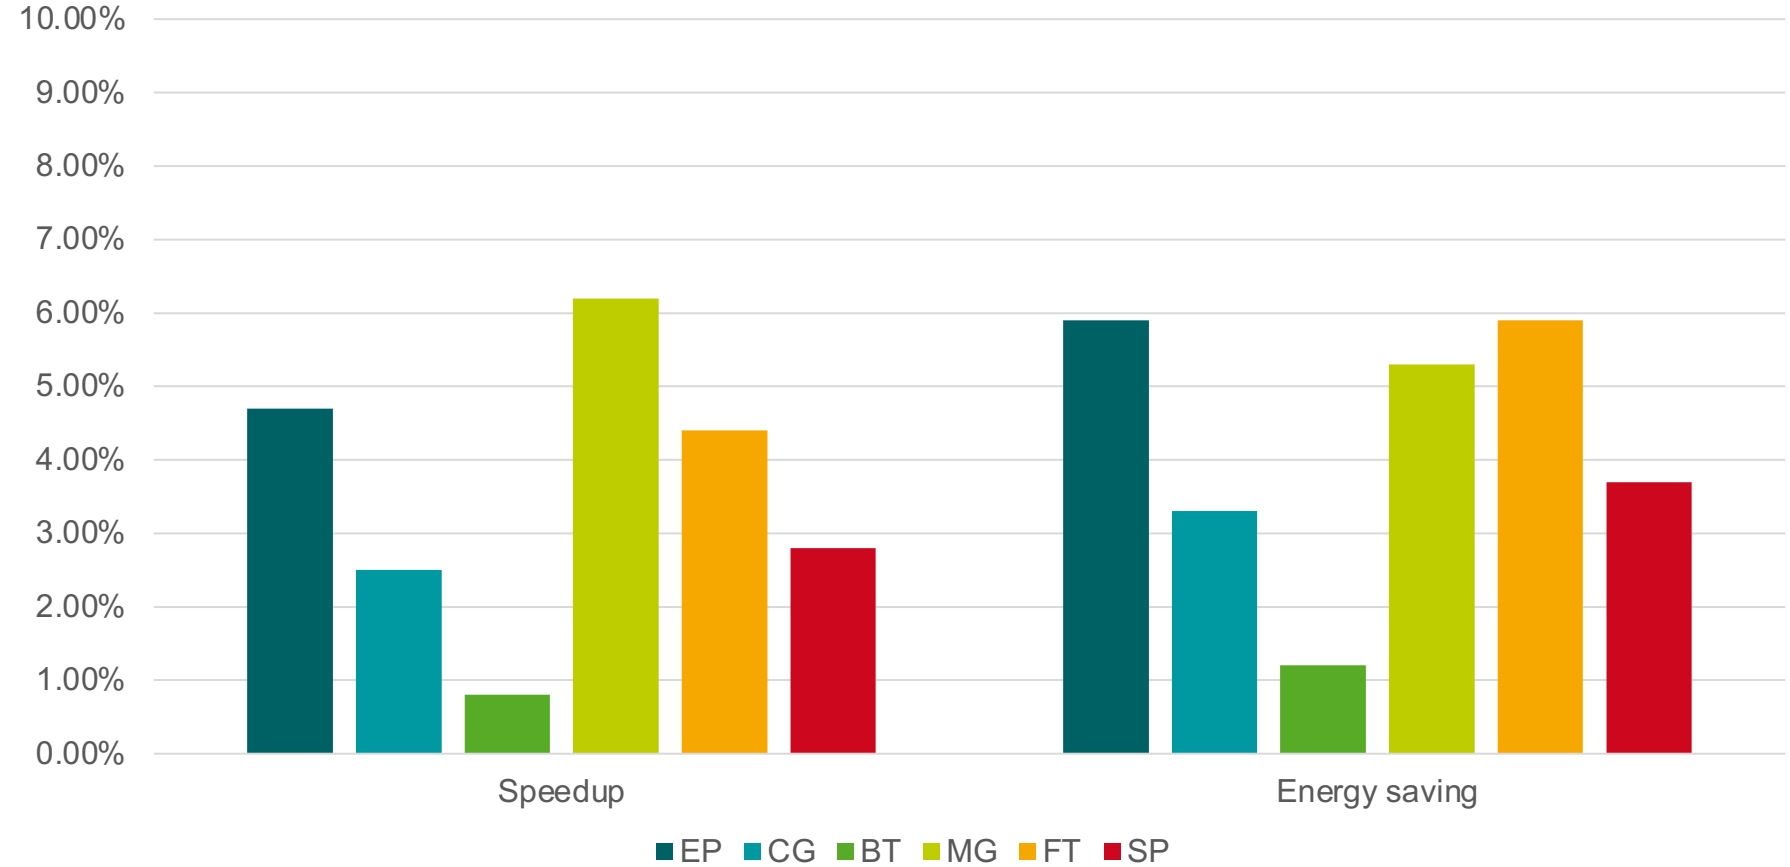
Slowdown under 4% for most benchmarks  
Worst case 6.5%

Up to 9.4% energy saving



### Speed up and energy saving [%] wrt. default under a power cap:69W

- Up to 6.2% speedup
- Up to 5.9% energy saving



# Summary

---

- Implement the Roofline model based on the sampling approach
- Extend the Roofline model with the time interval analysis method
  - Characterize different performance bottlenecks: compute-, memory- or transfer delay-bound
- The Extended Roofline achieves:
  - Without power cap:
    - Up to 13% package power saving and 9.4% energy saving
    - 2.5% ~ 6.5% slowdown
  - Under power cap:
    - Up to 6.2% speedup and 5.9% energy saving
- *Future Works:*
  - *Sampling by instructions*
    - *Sampling by instructions* triggers the interrupts only if there are program activities
    - More accurate and less overhead



---

# Thanks for your attention

## Citation

---

- [1] <https://www.top500.org/lists/top500/>
- [2] <https://techwireasia.com/2020/06/japans-fugaku-is-the-worlds-fastest-supercomputer/>
- [3] <https://crd.lbl.gov/assets/Uploads/ECP18-Roofline-1-intro.pdf>